
Quantile-Initialised Ordinal Collaborative Filtering with Gated Auxiliary Fusion

Kevin Zhu

Australian National University
u7338066@anu.edu.au

Abstract

We propose a simple calibration recipe for cumulative-link collaborative filtering: initialise the monotone thresholds from the empirical rating quantiles of the training set, parameterise their spacing through a softplus-inverse, and exclude them from L_2 regularisation. The recipe matches the data marginal exactly at the start of training and, on a transparent dot-product matrix-factorisation backbone, achieves low test negative log-likelihood (1.2584 across the five canonical MovieLens-100K splits, vs the $\log 5 \approx 1.61$ uniform baseline) without sacrificing RMSE. The recipe is paired with a per-dimension symmetric gated fusion of side information, which gives the largest RMSE change in our ablation, and a manifold-projection diagnostic that exposes item-side genre structure that linear projections miss (2-D IsoMap silhouette $+0.23$ vs $+0.05$ under PCA). The combined gated+ordinal configuration attains RMSE 0.9124 ± 0.0036 across all five MovieLens-100K splits (beating the no-fusion baseline 0.9160 and the additive-fusion variant 0.9189 on the across-split mean and on four of the five individual splits), with the same within-paper ordering observed on a $10\times$ -larger MovieLens-1M scale-up (RMSE 0.8481 ± 0.0006 , NLL 1.184). The contribution is practical rather than algorithmic: the components are standard; the recipe provides a calibrated probability output, with a small set of supporting diagnostics.

1 Introduction

Collaborative filtering (CF) for explicit feedback is commonly framed as predicting an unobserved scalar rating r_{ui} from a sparse user-item rating matrix $R \in \mathbb{R}^{n \times m}$ [17]. Latent-factor models based on matrix factorisation [8] remain standard baselines and building blocks on MovieLens [4]; later work includes neural hybrids [5, 21] and factorisation-machine variants [16, 3]. Two persistent issues remain when these models are applied as a course-style data-analytics tool rather than a production recommender. First, side information (user demographics, item genres) is typically injected by adding a projected feature vector to the latent embedding, with no learned mechanism to decide *when* to trust the embedding versus the side information. Second, the rating is treated as a continuous quantity even though it is in fact a five-class ordinal variable; the resulting predictions can be RMSE-competitive yet badly calibrated as distributions over star ratings.

Related work. The shortcomings above have each been addressed in isolation. *Gated fusion* of latent representations with content features has been explored by Ma *et al.*'s GATE [13], which uses a sigmoid gate to blend an item embedding with item content; gating is also implicit in attention-based factorisation [22]. *Cumulative-link ordinal regression* for recommendation has a long history starting with McCullagh's classical formulation [14], and was applied to MovieLens by Koren and Sill in OrdRec [9, 10]. Two recent papers operate in essentially the same setting as ours: OrdNMF [2] combines an ordinal head with non-negative matrix factorisation, and OPRFM [23] couples ordinal probit regression with factorisation machines on MovieLens-100K, 1M, and 10M. The cumulative-

link parameterisation itself has seen renewed interest in deep ordinal classification [20, 1] and in ordinal random fields for recommenders [12]. Less of this literature focuses on manifold visualisation of learned recommender embeddings; standard tools include IsoMap [18], t-SNE [19], and UMAP [15]. Our contribution is therefore not the invention of any single component, but their integration on a transparent linear core, together with a calibration recipe (empirical-quantile threshold initialisation) and a geometry diagnostic.

Contributions. We frame the contribution as implementation-level rather than algorithmic; each component below is itself well known.

1. *Headline.* A simple calibration recipe for cumulative-link CF: initialise the monotone thresholds from the training-set rating quantiles via $\theta_k = \text{logit}(\hat{F}_k)$, recover the unconstrained δ_j by inverting softplus, and exclude $\{\theta_1, \delta_j\}$ from L_2 regularisation. The predicted marginal then matches the data marginal at $s = 0$ to numerical precision and the ordinal NLL remains low throughout training.
2. *Supporting design choice.* A symmetric per-dimension gated fusion that blends ID embeddings with side-feature projections on both the user and the item sides, with zero-initialised gates so that all dimensions start from the same neutral 0.5 mixture. In our ablation this fusion accounts for the larger RMSE change — the ordinal head’s contribution is concentrated on calibration and per-class accuracy.
3. *Supporting diagnostic.* A manifold-projection lens (PCA, MDS, IsoMap) on the post-fusion embeddings, with silhouette scores stratified by occupation and dominant genre. IsoMap recovers item-side genre structure that linear projections miss (2-D silhouette $+0.23$ vs $+0.05$ under PCA), and the gap is itself an interpretable property of the linear core.

2 Problem Definition

Let $\mathcal{U} = \{1, \dots, n\}$ be the set of users and $\mathcal{I} = \{1, \dots, m\}$ the set of items. Observed ratings form a sparse set $\mathcal{D} = \{(u_t, i_t, r_t)\}_{t=1}^N$ with $r_t \in \{1, 2, 3, 4, 5\}$. Each user u carries a side-information vector $f_u \in \mathbb{R}^{d_u}$ (z-scored age, one-hot gender, one-hot occupation in MovieLens-100K, $d_u = 24$) and each item i carries $g_i \in \mathbb{R}^{d_i}$ (multi-hot genre indicator, $d_i = 19$). The task is to predict $\hat{r}_{ui}$ for held-out (u, i) pairs, evaluated by RMSE, MAE, classification accuracy of the rounded prediction, and where applicable the held-out negative log-likelihood (NLL) of the predicted rating distribution.

3 Method

3.1 Matrix factorisation backbone

We start from a biased dot-product MF model. Each user and item carries a d -dimensional embedding ($p_u, q_i \in \mathbb{R}^d$) and a scalar bias ($b_u, b_i \in \mathbb{R}$). The pre-head score is

$$s_{ui} = p'_u \cdot q'_i + b_u + b_i, \quad (1)$$

where p'_u, q'_i are the (possibly fused) representations defined below. With fusion disabled, $p'_u \equiv p_u$ and $q'_i \equiv q_i$.

3.2 Gated auxiliary fusion

Let $a_u = \text{ReLU}(W_u f_u + b_{f,u})$ be a learned projection of the user side features into the embedding space, and similarly $a_i = \text{ReLU}(W_i g_i + b_{f,i})$ for items. We define the fused user representation as

$$g_u = \sigma(W_g [p_u \parallel a_u] + b_g), \quad p'_u = (\mathbf{1} - g_u) \odot p_u + g_u \odot a_u, \quad (2)$$

where $\parallel$ denotes concatenation and $\odot$ is the Hadamard product, so the gate $g_u \in [0, 1]^d$ acts *per dimension*. The item-side fusion is identical with (q_i, a_i) in place of (p_u, a_u) . We zero-initialise both W_g and b_g , which makes $g_u \equiv \frac{1}{2} \mathbf{1}$ at training start. This gives every dimension the same neutral starting mixture between the ID embedding and the ReLU-projected side features; the gate then opens and closes per dimension as gradient flows. As an ablation we also consider the *additive* variant $p'_u = p_u + a_u$ and the *none* variant $p'_u \equiv p_u$ (and similarly for items).

3.3 Ordinal regression head with the calibration recipe

The headline contribution of this paper sits in this subsection: a small recipe for cumulative-link CF that holds the predicted rating distribution calibrated to the training-set marginal from step zero. The recipe has three ingredients — monotone-threshold parameterisation (Eq. 3–4), *empirical-quantile initialisation* (Eq. 5), and exclusion of the threshold parameters from L_2 regularisation. None of the three is individually new; their combination is what we contribute.

For the head we treat r as an ordinal variable with $K = 5$ classes. We use the cumulative-link model of McCullagh [14] with monotone thresholds $\theta_1 < \theta_2 < \theta_3 < \theta_4$:

$$P(r \leq k \mid s) = \sigma(\theta_k - s), \quad P(r = k \mid s) = \sigma(\theta_k - s) - \sigma(\theta_{k-1} - s), \quad (3)$$

with $\theta_0 = -\infty$ and $\theta_5 = +\infty$. We enforce monotonicity by the softplus reparameterisation

$$\theta_k = \theta_1 + \sum_{j=1}^{k-1} \text{softplus}(\delta_j), \quad (4)$$

which guarantees $\theta_k < \theta_{k+1}$ for any unconstrained $(\theta_1, \delta_1, \dots, \delta_{K-2})$. This trick has been used in deep ordinal classification [20, 1]. We use the logit link in (3); replacing the logistic CDF with the normal CDF (probit, as in OPRFM [23]) shifts headline RMSE by -0.0009 and NLL by $+0.005$ across the five splits, both within between-split standard deviation, so we treat the choice as a parameterisation detail rather than a methodological commitment.

Empirical-quantile initialisation (the recipe). Default zero initialisation (or a fixed $\theta_1 = -2$) does not match the training-set marginal and can start from a badly skewed class distribution. We instead initialise from the empirical rating distribution. Let $\hat{p}_k$ be the training-set frequency of rating k and $\hat{F}_k = \sum_{j \leq k} \hat{p}_j$ its cumulative. Choosing

$$\theta_k = \text{logit}(\hat{F}_k), \quad \delta_j = \log(\exp(\theta_{j+1} - \theta_j) - 1) \quad (5)$$

guarantees $P(r \leq k \mid s=0) = \hat{F}_k$, so the predicted marginal at the start of training equals the data marginal exactly. The second equation is just the softplus inverse, recovering the unconstrained δ_j that the optimiser then updates. Together with excluding $\{\theta_1, \delta_j\}$ from weight decay — their scale carries calibration information that L_2 shrinkage destroys — this is the calibration recipe we propose. It is ~ 10 lines of code, agnostic to the score s , and we use it on every ordinal cell in the experiments below.

The training loss is the negative log-likelihood $\mathcal{L}_{\text{ord}} = -\frac{1}{N} \sum_t \log P(r = r_t \mid s_{u_t t_t})$, and the point prediction is the expected rating $\hat{r}_{ui} = \sum_{k=1}^5 k P(r = k \mid s_{ui})$. As an ablation we also consider the standard *sigmoid* head: $\hat{r}_{ui} = 4\sigma(s_{ui}) + 1$ trained with mean-squared error.

Algorithm 1 Forward pass + ordinal NLL for one observed rating (u, i, r) .

- | | |
|---|---------------------------------------|
| 1: $a_u \leftarrow \text{ReLU}(W_u f_u + b_{f,u}); a_i \leftarrow \text{ReLU}(W_i g_i + b_{f,i})$
2: $g_u \leftarrow \sigma(W_g^{(u)}[E_u[u] \parallel a_u] + b_g^{(u)}); g_i \leftarrow \sigma(W_g^{(i)}[E_i[i] \parallel a_i] + b_g^{(i)})$
3: $p' \leftarrow (\mathbf{1} - g_u) \odot E_u[u] + g_u \odot a_u; q' \leftarrow (\mathbf{1} - g_i) \odot E_i[i] + g_i \odot a_i$
4: $s \leftarrow p' \cdot q' + b_u + b_i; \theta_k \leftarrow \theta_1 + \sum_{j < k} \text{softplus}(\delta_j)$
5: $P(r=k \mid s) \leftarrow \sigma(\theta_k - s) - \sigma(\theta_{k-1} - s)$, with $\theta_0 = -\infty, \theta_K = \infty$
6: return $\mathcal{L} = -\log P(r=r_{\text{true}} \mid s)$ | ▷ side projections
▷ per-dim gates |
|---|---------------------------------------|
-

3.4 Complexity

The MF backbone has $(n + m)d$ embedding parameters plus $n + m$ biases; for MovieLens-100K with $n = 943$, $m = 1682$, $d = 128$, that is $\approx 3.4 \times 10^5$ parameters. Each fusion module adds a projection ($d_{u/i} \cdot d + d$) and a gate ($2d^2 + d$); summed over both sides this is $\approx 7 \times 10^4$ parameters — about one-fifth of the embedding budget. The ordinal head adds $K - 1 = 4$ parameters. Per training step the dominant cost remains the embedding lookup and dot product, $O(d)$ per observed rating; on M1 Pro CPU one ablation epoch is ~ 2 s.

4 Experiments

4.1 Setup

We use MovieLens-100K [4] on its five canonical 80/20 splits `u1–u5`, each yielding 80,000 training and 20,000 test ratings on the same 943 users and 1,682 items. Headline numbers are reported as mean $\pm$ std across the five per-split means (each split-mean itself averages three training seeds), giving 15 trained models per ablation cell. The visualisation (§4.5) and cold-start (§4.6) experiments are restricted to the `u1` split. User side features are z-scored age, two-class one-hot gender, and 21-class one-hot occupation ($d_u = 24$); item features are 19-genre multi-hot ($d_i = 19$). All experiments use Adam [7] with learning rate 10^{-3} and weight decay $\lambda = 10^{-5}$ on all parameters except the ordinal-head thresholds. We carve a 10% random slice of `u1.base` (8,000 ratings) as a held-out validation set, leaving 72,000 for training; the val-split seed is fixed and shared across all training-seed replicates so val-split variance is not conflated with model-init variance. *Early stopping is driven by val-set RMSE only*, and the test set is read exactly once at the end of training, with the best-val weights restored. We report two protocols differing only in early-stopping aggressiveness: *headline* uses patience = 30 to allow the full 30-epoch budget, matching a fixed-budget training protocol of 30 epochs; *ablation* uses patience = 10 to keep the 18-run sweep tractable (~ 15 minutes wall on M1 Pro CPU). The two protocols agree on headline gated+ordinal RMSE to within ± 0.003 .

4.2 Baselines and metrics

We compare against three transparent baselines: (i) truncated SVD with rank 20 on the mean-centred rating matrix, (ii) biased MF trained with MSE, and (iii) NMF, identical to MF but with the user/item factor matrices constrained to be element-wise non-negative via projected gradient descent (clipping to ≥ 0 after each Adam step, biases unconstrained, following Lee & Seung [11]). Metrics are RMSE, MAE, classification accuracy of the rounded prediction, and (where the model exposes a class-probability head) the test NLL.

4.3 Results

Table 1 first places the proposed model against the three transparent baselines on the `u1` split under the headline protocol (single training seed). The stronger evidence is the multi-split ablation in Table 2, but this baseline table makes explicit that the model is not only being compared to its own ablated variants.

Table 1: Transparent baselines on MovieLens-100K `u1` under the headline protocol (seed 42, patience 30). The multi-split ablation in Table 2 is the main comparison. Lower is better for RMSE/MAE/NLL, higher for Accuracy.

Model	RMSE	MAE	Accuracy	NLL
SVD (rank 20)	1.0834	0.8896	0.3716	—
Biased MF	0.9221	0.7274	0.4124	—
NMF	0.9197	0.7252	0.4146	—
Gated+ordinal	0.9122	0.7130	0.4367	1.2472

The ablation matrix in Table 2 reports the full $\{\text{none, additive, gated}\} \times \{\text{sigmoid, ordinal}\}$ cross-product on three seeds. The proposed configuration — gated fusion with the ordinal head — has the best mean value on all four reported metrics. Across the five splits it improves RMSE by 0.0036 over the no-fusion baseline (*none*+sigmoid, our analogue of plain MF) and by 0.0065 over the additive-fusion variant under the sigmoid head. Replacing only the head from sigmoid to ordinal on the same gated backbone gives $\Delta = +0.0011$ RMSE (within between-split std), and improves MAE by 0.0056 and accuracy by 0.0087. The ordinal cell is also the only one that emits a native class distribution (NLL 1.2584). Because the RMSE differences are small relative to between-split variation, we interpret the main value of the ordinal head as probabilistic output and per-class behaviour rather than a large RMSE improvement.

Table 2: Ablation matrix on MovieLens-100K, mean $\pm$ std across the five canonical splits u1–u5 (each split-mean averages three seeds). Bold = best per column. Lower is better for RMSE/MAE/NLL, higher for Accuracy.

Fusion	Head	RMSE	MAE	Accuracy	NLL
none	sigmoid	0.9160 $\pm$ 0.0061	0.7241 $\pm$ 0.0059	0.4233 $\pm$ 0.0044	—
none	ordinal	0.9170 $\pm$ 0.0072	0.7204 $\pm$ 0.0069	0.4295 $\pm$ 0.0056	1.2634 $\pm$ 0.0050
additive	sigmoid	0.9189 $\pm$ 0.0044	0.7238 $\pm$ 0.0055	0.4243 $\pm$ 0.0057	—
additive	ordinal	0.9174 $\pm$ 0.0040	0.7173 $\pm$ 0.0035	0.4340 $\pm$ 0.0038	1.2654 $\pm$ 0.0017
gated	sigmoid	0.9135 $\pm$ 0.0037	0.7182 $\pm$ 0.0034	0.4286 $\pm$ 0.0054	—
gated	ordinal	0.9124 $\pm$ 0.0036	0.7126 $\pm$ 0.0042	0.4373 $\pm$ 0.0045	1.2584 $\pm$ 0.0043

4.4 Ablation study

Table 2 supports three claims.

Gated fusion gives the larger RMSE change. Holding the head fixed, replacing *none* fusion with *gated* drops RMSE by 0.0025 (sigmoid) or 0.0046 (ordinal); replacing *additive* with *gated* drops it by 0.0054 (sigmoid) or 0.0050 (ordinal). Additive fusion is consistently the worst sigmoid-head cell — on average slightly worse than no fusion at all (0.9189 vs 0.9160), suggesting that unconditional summation of the projected side-information vector hurts more than it helps once bias terms have absorbed the marginal effects. Per-dimension gating is the only fusion choice that improves the mean RMSE in this ablation.

The ordinal head buys native probabilities; its RMSE effect depends on fusion. Switching the head from sigmoid to ordinal moves RMSE by -0.0015 on additive fusion, -0.0011 on gated fusion, and $+0.0010$ on no fusion. All three are within between-split std. The sign pattern suggests that, in this experiment, ordinal prediction helps RMSE only when the backbone has informative inputs to discriminate across the five class boundaries; with no side information the predicted $P(r = k | s)$ collapses toward the marginal $\hat{p}_k$ and the ordinal NLL becomes a softer optimisation target than direct MSE on rounded ratings. What the ordinal head reliably buys — regardless of fusion — is a native rating distribution and per-class accuracy: it improves MAE by 0.003–0.007 and accuracy by 0.004–0.010 in every fusion regime, and it is the only native probabilistic head in the comparison.

Ordinal NLL is low across all ordinal cells. The proposed configuration achieves NLL 1.2584, well below the $\log 5 \approx 1.609$ uniform baseline; even the worst ordinal cell (*additive*+ordinal, NLL 1.2654) remains close to the proposed cell. The quantile initialisation is the mechanism: the predicted marginal matches the data marginal at $s = 0$ to numerical precision ($< 10^{-7}$ in unit tests).

Reliability diagram and per-class ECE. NLL is a single number; Figure 1 renders the underlying calibration curve directly. We compare the cumulative-link probabilities $P(r=k | s)$ of the gated+ordinal head against a Gaussian-kernel baseline that turns the gated+sigmoid point estimate $\hat{y}$ into a five-class distribution $P(r=k | \hat{y}) \propto \exp(-(k - \hat{y})^2 / 2\sigma^2)$ with $\sigma=1$ (so the sigmoid head can be evaluated under the same calibration definition). On pooled ECE the two heads are close (0.0108 kernelised sigmoid vs 0.0143 ordinal), but the per-class breakdown is the more informative read: the ordinal head’s ECE is lower on four of the five classes (class 1 0.009 vs 0.025, class 2 0.012 vs 0.036, class 3 0.019 vs 0.034, class 5 0.026 vs 0.036), with the kernelised sigmoid winning only on the dominant class 4 (0.009 vs 0.025). The pattern is the calibration version of the per-class F1 finding above: sigmoid+MSE concentrates probability mass on the dominant rating, which makes both its accuracy and its calibration look good there and worse everywhere else; the ordinal head spreads probability across the five classes more evenly.

Where the ordinal head wins: extreme classes. The aggregate accuracy and NLL numbers conceal a sharper mechanistic claim. Figure 2 reports per-class F1 on `u1.test` for the gated+sigmoid and gated+ordinal heads on the same backbone. The ordinal head wins on the boundary classes: $+0.098$ F1 on class 1 (a 54% relative improvement) and $+0.035$ F1 on class 5; classes 2–4 are within ± 0.015 of each other. Macro-F1 lifts from 0.341 (sigmoid) to 0.369 (ordinal), a +8% relative gain driven almost entirely by recall on classes 1 and 5 (class-1 recall 0.169 vs 0.102; class-5




Figure 1: Reliability diagrams on `u1.test` for the gated backbone under each head. Predicted probability is bucketed into 10 bins and plotted against observed empirical frequency; perfect calibration sits on the diagonal. The faint per-class lines reveal the structure of the calibration error that pooled ECE conceals: gated+ordinal is better calibrated on classes 1, 2, 3, and 5; gated+sigmoid (Gaussian kernelised) wins only on the dominant class 4.

recall 0.235 vs 0.204). The sigmoid head trained with MSE biases predictions toward the middle of the rating range; the cumulative-link parameterisation explicitly carries thresholds for each class boundary, and uses them to recover ratings the sigmoid head rounds back into class 3 or class 4. This is the per-class footprint of the calibration claim, not just an aggregate.




Figure 2: Per-class F1 on `u1.test` for the same gated backbone with sigmoid vs. ordinal heads. The ordinal head wins on the extreme classes (1 and 5) where the cumulative-link thresholds carry boundary information; classes 2–4 are within seed-level noise.

Figure 3 shows the sensitivity sweep over the embedding dimension d and the weight-decay coefficient λ , with all other settings held at the defaults of Section 4. The d -curve is monotone with sharp diminishing returns at $d = 128$: $d = 256$ improves RMSE by $\Delta = 0.0024$ at +34% wall time, comparable to the seed-to-seed variance at $d = 128$ (0.0029), which we treat as too small a gain to justify the compute cost. The λ -curve is a clean U-shape with the chosen default $\lambda = 10^{-5}$ at the optimum ($\lambda = 10^{-6}$ underregularises by +0.0018, $\lambda = 10^{-4}$ overregularises by +0.0080). We did not tune hyperparameters post hoc.




Figure 3: Sensitivity of test RMSE to embedding dimension d (left) and weight decay λ (right) on the gated+ordinal configuration. Error bars are ± 1 std over three seeds.

4.5 Embedding visualisation

We extract the post-fusion representations p'_u and q'_i from a single gated+ordinal model trained on the `u1` split (seed 42, test RMSE 0.9165), and project them to $\mathbb{R}^2$ via three classical manifold methods: PCA [6] (linear, variance preserving), classical MDS (linear, distance preserving), and IsoMap [18] (nonlinear, geodesic distance preserving, $k = 15$ neighbours). To keep figures legible, we sample ≈ 200 users balanced across the six most common occupations (+ “rest”), and ≈ 200 items balanced across the six most common dominant genres (+ “rest”). Cluster quality is measured by the silhouette score with respect to the categorical labels.




Figure 4: Two-dimensional projections of the post-fusion user representations p'_u (top, coloured by occupation) and item representations q'_i (bottom, coloured by dominant genre). Columns show PCA, classical MDS, and IsoMap with $k = 15$ neighbours. Item-side genre structure is recovered most clearly by IsoMap; user-side occupation structure is present in the high-dimensional space but does not survive projection to two dimensions.

Two observations stand out (Table 3, Figure 4). First, both entities carry positive high-dimensional silhouette: the model has, without any auxiliary classification objective, learned representations

Table 3: Silhouette scores in the full 128-dimensional space (HD) and after projection to two dimensions, with respect to occupation labels (users) and dominant genre labels (items).

	HD ($d = 128$)	2D PCA	2D MDS	2D IsoMap
p'_u by occupation	+0.0848	−0.107	−0.103	−0.114
q'_i by dominant genre	+0.1930	+0.048	+0.137	+0.232

that are partially side-information aligned. Item-side genre structure is roughly twice as strong as user-side occupation structure (+0.193 vs +0.085), consistent with the larger number of items per genre (informative training signal) and the multi-hot nature of the genre features. Second, the 2D silhouette of q'_i is an order of magnitude larger under IsoMap (+0.232) than under PCA (+0.048): the genre manifold is curved, and a linear projection collapses points from different genres on top of one another. This is the diagnostic that classical PCA-based recsys visualisations miss. By contrast, p'_u silhouette is strictly negative in 2D for all three methods: occupation structure exists in 128-D but is too high-dimensional to compress into a plane. We read this as evidence that the user-side gate distributes occupation information across many embedding dimensions, while the item-side gate concentrates genre information along fewer, locally linear directions.

4.6 Gate behaviour and cold-start regime

Per-user trajectory. A standard intuition for a learned gate is that it should *close* as the ID embedding accumulates training signal — a power user with hundreds of observed ratings should rely on the ID embedding alone, while a cold user with a handful of ratings should rely on the side-feature projection. We test this by stratifying every test prediction by the corresponding user’s training activity $|R_u|$ into five activity buckets, and measuring the per-bucket mean gate value and per-bucket RMSE for both the gated+ordinal and additive+ordinal models (Figure 5, Table 4). The activity buckets contain {189, 108, 89, 59, 14} users covering {2717, 3441, 5846, 6350, 1646} test ratings.




Figure 5: Per-bucket cold-start analysis. Left: mean per-prediction gate value across user-activity buckets. The user gate g_u rises modestly with $|R_u|$ from ≈ 0.25 to ≈ 0.29 , while g_i stays essentially flat near 0.32; both remain well below the zero-init value of 0.5. Right: test RMSE by user-activity bucket for the two fusion variants. Gated fusion’s advantage peaks in the mid-tail and flips sign on the power-user tail, where additive narrowly wins.

Table 4: Cold-start RMSE breakdown by user activity. Negative Δ means additive fusion wins on that bucket.

$ R_u $ bucket	< 30	30–59	60–119	120–239	≥ 240
gated RMSE	0.9799	0.9649	0.8841	0.8896	0.9200
additive RMSE	0.9803	0.9768	0.8929	0.8946	0.9128
Δ (add. – gated)	+0.0004	+0.0119	+0.0088	+0.0051	−0.0071
mean g_u	0.246	0.256	0.263	0.268	0.291
mean g_i	0.316	0.317	0.316	0.317	0.316

The data do not support the naive intuition in this setting. First, g_u moves modestly with $|R_u|$, increasing from 0.246 on the cold tail to 0.291 on power users (an 18% rise), while g_i stays essen-

tially constant at 0.316–0.317. The user gate *opens* as the user accumulates ratings — the opposite of the naive “gate closes with data” intuition. Both gates remain well below the zero-init value of 0.5, so the ID-embedding pathway has larger weight throughout, but the relative weight given to the side-feature projection grows on the user side as more training signal becomes available. Second, the gated+ordinal advantage over additive+ordinal peaks in the mid-activity bucket (30–59 ratings, $\Delta = +0.012$ RMSE), shrinks monotonically through the middle of the population, ties on the cold tail (<30 ratings, $\Delta = +0.0004$, within seed variance), and *flips sign* on the power-user tail (≥ 240 ratings, $\Delta = -0.007$, i.e. additive narrowly wins). For users with hundreds of training ratings, the ID embedding already carries enough signal that mixing in the side-info projection through the gate appears to introduce variance without added information — and the gate’s continued opening at the high-activity end of the population may be part of why. We read this as evidence that learned gating is not uniformly beneficial across user-activity regimes in this linear-core CF setting. This complements related gated and probabilistic recommender work (GATE [13], OPRFM [23]) by showing that aggregate gains should not automatically be read as uniform across user activity regimes. The practical implication is that gated fusion is most valuable on mid-tail users; for the power-user tail, an additive (or no-fusion) variant may be preferable, though this needs confirmation beyond the single u1 cold-start analysis.

4.7 Per-dimension gate interpretability

Aggregate gate values disguise a striking per-dimension pattern. For each of the $d=128$ embedding dimensions of the trained gated+ordinal model (u1, seed 42), Figure 6 sorts dimensions by their population-mean gate value and shades ± 1 standard deviation around the mean. The user-side gate g_u spans $[0.002, 0.984]$ across dimensions with population-mean 0.263; the item-side gate g_i spans $[0.002, 0.992]$ with population-mean 0.317. Only 20.3% of user-side dimensions and 21.9% of item-side dimensions sit above the zero-init value of 0.5. The distribution is therefore strongly bimodal: most dimensions *specialise* to the ID-embedding pathway (gate near zero) while a smaller subset specialises to the side-information projection (gate near one). The fusion mechanism is dividing labour across dimensions rather than averaging the two sources.




Figure 6: Per-dimension gate distribution on `u1.test`, sorted by mean. The shaded band is ± 1 std across the test population. Both sides exhibit a long left tail of fully closed dimensions (gate $\rightarrow 0$, ID embedding only) and a right tail of fully open dimensions (gate $\rightarrow 1$, side-info projection only); the fusion mechanism specialises rather than averages.

Is the gate’s choice population-wide or category-conditional? We stratify the per-prediction gate by the user’s occupation (top six categories) and by the item’s dominant genre (top six), then compute the pairwise Spearman rank correlation between the per-dimension orderings of any two strata. For occupations the mean rank correlation is 0.96; for genres it is 0.91. Two strata picking similar dimensions to open implies the gate is making a population-wide decision rather than conditioning on the categorical side information. So while the *fused output* of course depends on the user’s and item’s features (the side-information vector is plumbed into both the projection and the gate input), the *per-dimension partition* between ID-driven and side-info-driven dimensions is a population-level architectural choice that emerges from training.

4.8 Scale-up to MovieLens-1M

To check whether the proposed integration remains competitive on a $10\times$ larger benchmark, we run the three most informative cells — the no-fusion baseline (*none*+sigmoid), gated fusion alone (*gated*+sigmoid), and the proposed gated+ordinal — on MovieLens-1M [4] (1,000,209 ratings, 6,040 users, 3,706 rated items; $d_u = 30$ from gender, 7-bucket age and 21-class occupation; $d_i = 18$ from a multi-hot genre vocabulary derived from the data). MovieLens-1M does not ship canonical CV splits, so we use a deterministic random 80/10/10 train/val/test partition keyed by a fixed split seed; all other settings (Adam, $d = 128$, $\lambda = 10^{-5}$, ablation patience 10) match the MovieLens-100K experiments. Each cell is averaged over the same three training seeds.

Table 5: Scale-up to MovieLens-1M (random 80/10/10 split, mean $\pm$ std over three seeds). Bold = best per column. The proposed configuration retains the best mean metrics among the three tested cells on a $10\times$ larger benchmark, with substantially tighter seed variance.

Fusion	Head	RMSE	MAE	Accuracy	NLL
none	sigmoid	0.8567 ± 0.0008	0.6782 ± 0.0007	0.4480 ± 0.0005	—
gated	sigmoid	0.8502 ± 0.0003	0.6695 ± 0.0001	0.4579 ± 0.0002	—
gated	ordinal	0.8481 ± 0.0006	0.6630 ± 0.0010	0.4654 ± 0.0014	1.1842 ± 0.0013

Two observations are worth flagging. First, the ordering of the three tested cells matches the MovieLens-100K result (Table 2): gated fusion lifts RMSE over no-fusion by 0.0086 (vs 0.0025 on 100K), and adding the ordinal head on top contributes a further 0.0021 (vs 0.0011 on 100K, also within seed noise). The integration story carries unchanged to the larger benchmark, with the fusion’s contribution roughly tripling in absolute magnitude. Second, the absolute RMSE drops by ~ 0.064 between the two datasets (0.9124 on the five-split 100K mean to 0.8481 on 1M); the gated+ordinal cell produces an ordinal NLL of 1.184, substantially below the corresponding log 5 uniform baseline. Seed variance is roughly $5\times$ tighter on 1M ($\text{std} \approx 10^{-3}$ versus $\sim 4 \times 10^{-3}$ on 100K), consistent with the larger training set’s stabilising effect on the gate. Per-cell wall time scaled roughly linearly: a single ablation-protocol gated+ordinal run takes ~ 12 minutes on M1 Pro CPU compared to ~ 50 seconds on 100K.

4.9 Discussion

The ablation paints a sharper picture than the headline number alone. The fusion mechanism accounts for the larger part of the empirical RMSE change: gated fusion lifts RMSE on the across-split mean under both heads (and on four of the five individual splits; u4 is the lone exception, where the gated configuration is within seed std of the no-fusion baseline). The ordinal head provides native probability estimates cheaply — four learnable threshold parameters — but does not help RMSE on its own; only when the backbone has informative inputs to discriminate against does the cumulative-link parameterisation pay off in our runs. This interaction suggests that ordinal heads are most useful here when paired with side-information mechanisms, not as a stand-alone substitute for them.

Limitations. The benchmark is small (100K ratings); the side information is limited to demographics and genres; and the ordinal head is calibrated to the rating-class marginal but does not model per-user heteroscedasticity.

5 Conclusion

We presented a calibration recipe for cumulative-link collaborative filtering: initialise the monotone thresholds from the training-set rating quantiles, parameterise their spacing through a softplus-inverse, and exclude them from L_2 regularisation. The recipe is implementation-level (~ 10 lines of code, agnostic to the score function), and it produces a head whose predicted marginal matches the data marginal at training start to numerical precision and whose ordinal NLL remains low throughout training. Paired with a per-dimension symmetric gated fusion of side information, the resulting model attains RMSE 0.9124 ± 0.0036 and NLL 1.2584 across the five canonical MovieLens-100K splits under val-driven early stopping, improving the across-split mean over both an MF baseline

and an additive-fusion variant on the main reported metrics. The ablation isolates fusion as the headline mechanism and reveals that the ordinal head’s benefit is conditioned on the presence of informative inputs, an interaction that matters when deciding whether ordinal regression is worth its small parameter cost. Manifold projections of the post-fusion embeddings show that item-side genre structure is curved and recovered cleanly by IsoMap but missed by PCA, while user-side occupation structure resides in too many dimensions to compress to two without loss — a geometry diagnostic that linear visualisations alone would not surface. A cold-start stratification reveals that the user-side gate *opens* modestly as users accumulate ratings (counter to the naive “gate closes with data” intuition), and that gated fusion’s advantage over additive fusion is a mid-tail phenomenon: it ties on the cold tail and reverses sign on the power-user tail, suggesting that learned-gating mechanisms in linear-core CF may need an additive or no-fusion fallback at the activity extremes. Repeating the three most informative cells on MovieLens-1M (10× larger, 1 000 209 ratings) gives the same within-paper ordering and substantially tighter variance: gated+ordinal attains RMSE 0.8481 ± 0.0006 and NLL 1.1842, the no-fusion baseline 0.8567 ± 0.0008 .

References

- [1] Wenzhi Cao, Vahid Mirjalili, and Sebastian Raschka. Rank consistent ordinal regression for neural networks with application to age estimation. *Pattern Recognition Letters*, 140:325–331, 2020.
- [2] Olivier Gouvert, Thomas Oberlin, and Cédric Févotte. Ordinal non-negative matrix factorization for recommendation. In *ICML*, pages 3680–3689, 2020.
- [3] Huifeng Guo, Ruiming Tang, Yunming Ye, Zhenguo Li, and Xiuqiang He. DeepFM: a factorization-machine based neural network for CTR prediction. In *IJCAI*, pages 1725–1731, 2017.
- [4] F Maxwell Harper and Joseph A Konstan. The MovieLens datasets: History and context. *TiiS*, 5(4):1–19, 2015.
- [5] Xiangnan He, Lizi Liao, Hanwang Zhang, Liqiang Nie, Xia Hu, and Tat-Seng Chua. Neural collaborative filtering. In *WWW*, pages 173–182, 2017.
- [6] Ian T Jolliffe. *Principal Component Analysis*. Springer, 2002.
- [7] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv:1412.6980*, 2014.
- [8] Yehuda Koren, Robert Bell, and Chris Volinsky. Matrix factorization techniques for recommender systems. *Computer*, 42(8):30–37, 2009.
- [9] Yehuda Koren and Joseph Sill. OrdRec: an ordinal model for predicting personalized item rating distributions. In *RecSys*, pages 117–124, 2011.
- [10] Yehuda Koren and Joseph Sill. Collaborative filtering on ordinal user feedback. In *IJCAI*, pages 3022–3026, 2013.
- [11] Daniel Lee and H Sebastian Seung. Algorithms for non-negative matrix factorization. *Advances in Neural Information Processing Systems*, 13, 2000.
- [12] Shaowu Liu, Truyen Tran, and Gang Li. Ordinal random fields for recommender systems. In *Proceedings of the Sixth Asian Conference on Machine Learning (ACML)*, volume 39 of *Proceedings of Machine Learning Research*, pages 283–298. PMLR, 2015.
- [13] Chen Ma, Peng Kang, Bin Wu, Qinglong Wang, and Xue Liu. Gated attentive-autoencoder for content-aware recommendation. In *WSDM*, pages 519–527, 2019.
- [14] Peter McCullagh. Regression models for ordinal data. *JRSS B*, 42(2):109–127, 1980.
- [15] Leland McInnes, John Healy, and James Melville. UMAP: Uniform manifold approximation and projection for dimension reduction. *arXiv:1802.03426*, 2018.
- [16] Steffen Rendle. Factorization machines. In *ICDM*, pages 995–1000, 2010.
- [17] J Ben Schafer, Dan Frankowski, Jon Herlocker, and Shilad Sen. Collaborative filtering recommender systems. In *The adaptive web*, pages 291–324. Springer, 2007.
- [18] Joshua B Tenenbaum, Vin De Silva, and John C Langford. A global geometric framework for nonlinear dimensionality reduction. *Science*, 290(5500):2319–2323, 2000.

- [19] Laurens van der Maaten and Geoffrey Hinton. Visualizing data using t-SNE. *JMLR*, 9:2579–2605, 2008.
- [20] Víctor Manuel Vargas, Pedro Antonio Gutiérrez, and César Hervás-Martínez. Cumulative link models for deep ordinal classification. *Neurocomputing*, 401:48–58, 2020.
- [21] Hao Wang, Naiyan Wang, and Dit-Yan Yeung. Collaborative deep learning for recommender systems. In *KDD*, pages 1235–1244, 2015.
- [22] Jun Xiao, Hao Ye, Xiangnan He, Hanwang Zhang, Fei Wu, and Tat-Seng Chua. Attentional factorization machines: learning the weight of feature interactions via attention networks. In *IJCAI*, pages 3119–3125, 2017.
- [23] Nilufar Zaman and Angshuman Jana. Automated recommendation model using ordinal probit regression factorization machines. *International Journal of Data Science and Analytics*, 20(3):2615–2629, 2025.

A Reproducibility

All experiments run on a 14-inch MacBook Pro M1 Pro CPU (no GPU) under Python 3.11 with PyTorch 2.11 (CPU build), NumPy 2.4, scikit-learn 1.8 and matplotlib 3.10. The code is structured as a small, dependency-light package; the entire test suite (six unit tests covering gate shape, zero-init invariant, ordinal probability normalisation, threshold monotonicity, NMF non-negativity projection, and end-to-end forward/backward) runs in ~ 2 s. Every reported headline number traces to a JSON dump in the repository archive.

Table 6: Reproducibility commands. All scripts are deterministic given the listed seeds; `val_seed=0` is fixed across the project.

Result	Command	Wall time
6-cell ablation, u1, 3 seeds	<code>run_ablations.py</code>	~ 10 min
Sensitivity sweep (d, λ), 3 seeds	<code>run_sensitivity.py</code>	~ 15 min
5-split CV, 6 cells, 3 seeds (90 runs)	<code>run_multisplit.py</code>	~ 60 min
Manifold visualisation (u1)	<code>run_visualization.py</code>	~ 1 min
Cold-start regime analysis (u1)	<code>run_coldstart.py</code>	~ 2 min
Per-class confusion / F1 (u1)	<code>run_classmetrics.py</code>	~ 2 min
Logit vs. probit link (5 splits)	<code>run_link_compare.py</code>	~ 10 min
Per-dim gate analysis (u1)	<code>run_gate_analysis.py</code>	~ 1 min
Reliability diagram + ECE (u1)	<code>run_reliability.py</code>	~ 2 min
MovieLens-1M scale-up (3 cells $\times$ 3 seeds)	<code>run_mllm.py</code>	~ 85 min

Seeds: training-seed set $\{42, 43, 44\}$; val-split seed 0 (independent of training). Optimiser Adam [7] with learning rate 10^{-3} and weight decay $\lambda = 10^{-5}$; ordinal-head thresholds excluded from weight decay. Patience= 10 for the ablation protocol and 30 for the headline protocol. Embedding dimension $d = 128$, batch size 64, max 30 epochs. NMF baseline uses projected gradient descent on the factor matrices (clamping to ≥ 0 after each Adam step), with biases unconstrained.

B Software

The repository includes a single `uv` lockfile, the deterministic dataset loader, all experiment scripts, the six-test `pytest` suite, and a per-result JSON archive that includes the configuration snapshot used to produce it. Every figure and number in this paper can be regenerated by running the corresponding script in Table 6.